

HW1

Q1

I am surprised about how optimistic it all started out. It was just a proposal of a few scientists getting together during summer and people thought it would be done by then. I was also surprised by the perceptron, how there was a mass amount of physical wires that simulated a neural network. I never thought a neural network could be a physical thing. I wasn't really sure how the insides of chatgpt worked, so I was surprised by the number of transformer blocks, but it makes sense.

I'm not surprised that AI still has a lot of faults, like cars not recognizing that the people are actually a part of an ad, or recognizing images at weird angles. I'm also not surprised that LLMs can't generalize problems, like if you change a problem slightly it will most likely get it wrong.

For the future, I agree that AI will free us from tedious and dangerous jobs, but I also think it will take away a lot of jobs from people. Either forcing people to be jobless or pursue other fields. I also don't think AI tools will help much in creativity. AI will definitely create faster and new things we haven't seen before, but I still think humans creating art or music or film is more natural. It is very hard to stop people from abusing the power of AI, so the pros and cons are very hard to balance. To be honest I'm not really leaning towards either side about where it will go in the future, whether it will continue the surge or drop to another "AI winter".

Q2

I liked the counter argument of how raising the stakes doesn't make people judge more correctly, because my initial thought was also that the statement would hold true. So I was surprised that the examples he provided were sound. I was also surprised how picking in the NFL could be modeled like this. We can really put a concrete model on basically everything. It was interesting seeing all the examples of irrelevant facts in our lives, how something could affect another thing so much, but are in reality irrelevant.

I have definitely encountered "availability heuristics" happen before, where people tend to judge something based on how big of a news they hear online, like plane crashes versus car crashes. I've heard of the "no free market" hypothesis before, about how you can't beat the market, and I think it's very true how gaining one thing requires one to sacrifice another.

I think the main thing was how humans are "dumber" and "nicer" than econs, which makes us human. And I think that's also the line between humans and AI, and how far between that line we want to go. He also talked about how it is easier to make money by catering to consumer's biases than correcting them. The way we train AI is to make the computer system act more like humans, adjust to the biases of humans. The bias of humans is what makes it "right" sometimes, and we keep training and feeding information into LLMs that reflect what we as humans would do.

HW2

Q1

I was initially confused why the emotional disorder where some areas of the brain is damaged so that the patient is super unemotional had anything to do with artificial intelligence and emotional intelligence. But then I thought about how being unemotional, and thus being able to solve everything logically, is much like a computer. If we think about it, there is a huge amount of emotional factor that goes into decision making each day, and how much it matters to go down deeper into a rabbit hole or just make a decision.

I agree that the people who rise to the top aren't those who are exceptionally smart, but people who understand people. The best combination for humans would be a balance of smarts and emotional intelligence. However, I think that AI should primarily be focused on being academically smart, and not trying to be more human. Besides the fact that it would be extremely hard to train AI to better recognize emotion, because of context cues and simply how we are as humans, I don't think it would be a very useful use of our time and effort to do so.

It does surprise me that people are trying to instill more emotional intelligence on AI and robots, especially chat bots. I agree that when someone encounters a problem, most times they want to be understood first before given a solution. But solving loneliness and how people become attached to AI chatbots makes me uncomfortable to think about.

Q2

Dario Amodei stresses that AI shouldn't avoid controversial topics, since ignoring them could mean missing out on important discussions about risks and ethics. He talks about ethical traps—situations where AI could lead to unexpected problems, making it even more important to build in safety measures. He says that AI would be able to do things end to end soon, like searching up reservations then booking them, and giving feedback / suggestions. I think we're already seeing that in chatGPT, but I don't think AI should go much further than helping us book and plan events.

In the video, Dario Amodei talked about the one shot persuasion, where he put a group of humans versus AI to write a persuasive essay. I think it is very scary how we can feed AI human data, have it analyze everyone's weak spots, set it loose, and have it able to persuade mostly anyone. This creates the concern that AI might be able to create a religion some day, where its followers will believe everything it says.

What surprised me was that he thinks that some things will become clearer very soon, but I think we're still very confused about where AI will go, and how it might be abused. I think this challenges the common perception that such developments are decades away. His confidence in the scalability of current AI models to achieve near-human intelligence diminishes the urgency for robust safety measures and ethical guidelines. The first podcast was more uncertain and vague, while in the second video, Dario Amodei is portrayed more certain of AI's potential.



Q1

When we have smaller group sizes, for around 85% accuracy, we need around 50% of the individuals to be informed. However, for larger groups and the same accuracy, we only need less than 5% to be informed. I think this is very true because once there is a large enough group, a sense of identity is lost. Whether it is humans or animals, all it might take is those couple individuals to start something, then soon everyone, seeing everything else, will follow.


Q2

When a super opinionated minority takes over, they can win almost every time, like 98% of the time. But if you throw in just 10 or 12 uninformed people, everything shifts, and the majority regains control. Basically, those neutral people break the hold of the strong opinions and make the decision-making process more fair. It's kind of wild because you'd think uninformed people wouldn't matter much, but they actually keep things from being hijacked by a loud minority.

Swarm



Torus



Parallel Group

